



Welcome to

Network Automation and Monitoring Basics

How to manage networks

Henrik Kramselund Jereminsen hkj@zencurity.com @kramse  

Slides are available as PDF, kramse@Github
`network-automation-monitoring.tex` in the repo `security-courses`

Contact information



- Henrik Kramselund Jereminsen, internet samurai mostly networks and infosec
- Independent network and security consultant
- Master from the Computer Science Department at the University of Copenhagen, DIKU
- Email: hkj@zencurity.dk Mobile: +45 2026 6000

You are welcome to drop me an email

Goals: Network Automation and Monitoring Basics



- Introduce network management suitable for modern times
- Introduce resources, programs, people, authors, documents, sites that further your exploration into network management
- Starting with manual tools, move to some excellent automated ones
- I recommend open source tools, feel free to try these and then decide to pay for commercial ones, if you like



Plan for today

A starter pack for network management

- Network information TCP/IP
- The basic tools for diagnosing network problems
- Smokeping for automated monitoring
- LibreNMS - a fully featured network monitoring system
- Ansible for robust configuration of servers and other devices
- Oxidized gathering network device configuration
- Syslog server - Logstash, Elasticsearch, Dashboards

Duration: 4 hours - with breaks

Keywords: IP address plans, core infrastructure, SNMP, DNS, DHCPD, Netflow, dashboards, TCP, UDP, ICMP, routing, switching, a little BGP, Ansible for network devices with Junos, OpenBSD and Linux as examples, Netflow, sFlow, monitoring systems, LibreNMS, Oxidized, Smokeping

Time schedule



- 17:00 - 18:15
Introduction and basic manual tools
- 30min break
- 18:45 - 19:30 45min
Recommended tools, Nipap, Smokeping, LibreNMS
- 15min break
- 19:45 - 21:00
Centralized solutions, tie together
with break somewhere

About equipment and exercises



- Bringing a laptop to my courses is not required, but welcome
- Links etc. are in the slides and open source licensed, PDFs
- Exercises booklets are available for many of my courses, see Github
but it is expected that participants will do any exercises on their own later or at the scheduled hacker days
- The hacker days will be announced in various places
- Events like BornHack are excellent places to arrange hacker days in the network warrior village, or other places

Invite a few friends, make a hacker day and work together!

Course Materials



This material is in multiple parts:

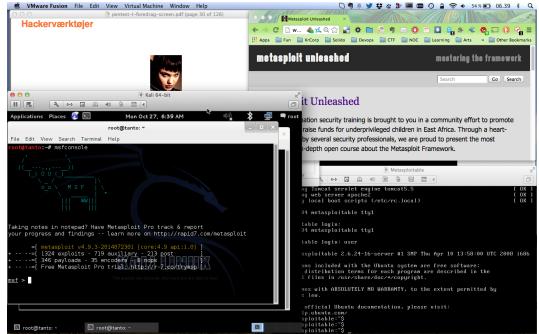
- Slide shows - presentation - this file
- Exercises - PDF files in my repository

Links

- All materials will be released as open source at:
<https://github.com/kramse/security-courses/>
- Additional resources from the internet linked from lecture plans:
<https://zencurity.gitbook.io/kea-it-sikkerhed/>

Note: slides and materials will mostly be in english, but presentation language will be danish

Hackerlab Setup



- Hardware: modern laptop CPU with virtualisation
Dont forget to enable hardware virtualisation in the BIOS
- Software Host OS: Windows, Mac, Linux
- Virtualisation software: VMware, Virtual box, HyperV pick your poison
- Hackersoftware: Kali Virtual Machine <https://www.kali.org/>



Networking Hardware

If you want to do exercises and work with networks

It will be an advantage to have a wireless USB network card and an USB Ethernet card.

The following are two models I have used:

- TP-link TL-WN722N hardware version 2.0 cheap
- Alfa AWUS036ACH 2.4GHz + 5GHz Dual-Band and high performing
- Often you need to compile drivers yourself, and research a bit

Getting an USB card allows you to use the regular one for the main OS, and insert the USB into the virtual machine

Network Management

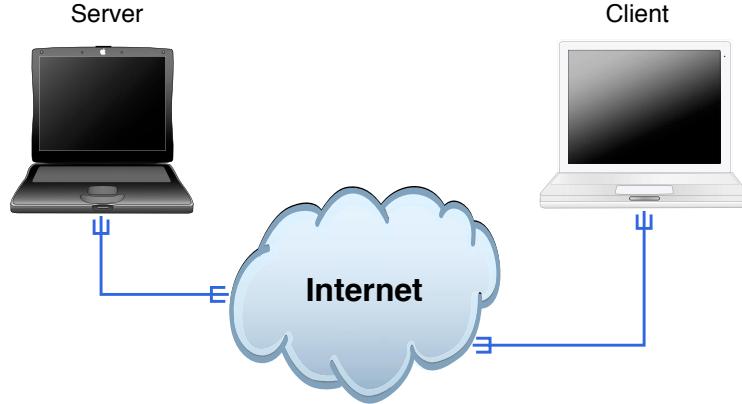


Network management is the process of administering and managing computer networks. Services provided by this discipline include fault analysis, performance management, provisioning of networks and maintaining the quality of service. Software that enables network administrators to perform their functions is called network management software.

https://en.wikipedia.org/wiki/Network_management

- What are we talking about today
- Complex modern networks
- Hint: easier to consider your network a critical resource, and start monitoring now

Internet Today



- Clients and servers, roots in the academic world
- Protocols are old, some more than 20 years
- Very little is encrypted, mostly HTTPS

Internet is Open Standards!



We reject kings, presidents, and voting.
We believe in rough consensus and running code.
– The IETF credo Dave Clark, 1992.

Request for comments - RFC - er en serie af dokumenter

RFC, BCP, FYI, informational

de første stammer tilbage fra 1969

Ændres ikke, men får status Obsoleted når der udkommer en nyere version af en standard

Standards track:

Proposed Standard → Draft Standard → Standard

Åbne standarder = åbenhed, ikke garanti for sikkerhed

What is the Internet



Communication between humans - currently!

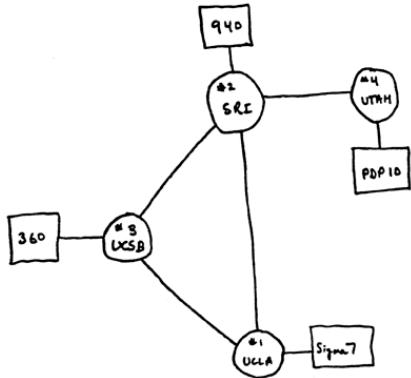
Based on TCP/IP

- best effort
- packet switching (IPv6 calls it packets, not datagram)
- *connection-oriented* TCP
- *connection-less* UDP

RFC-1958:

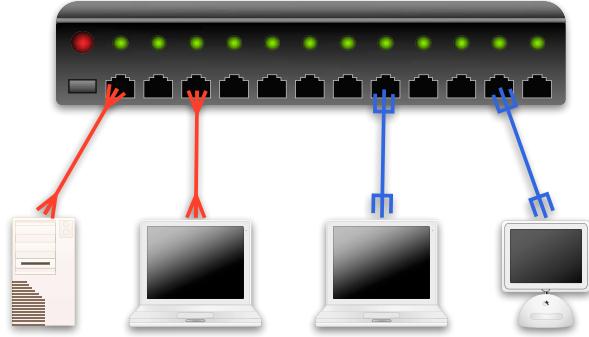
A good analogy for the development of the Internet is that of constantly renewing the individual streets and buildings of a city, rather than razing the city and rebuilding it. The architectural principles therefore aim to provide a framework for creating cooperation and standards, as a small "spanning set" of rules that generates a large, varied and evolving space of technology.

Internet historisk set - anno 1969



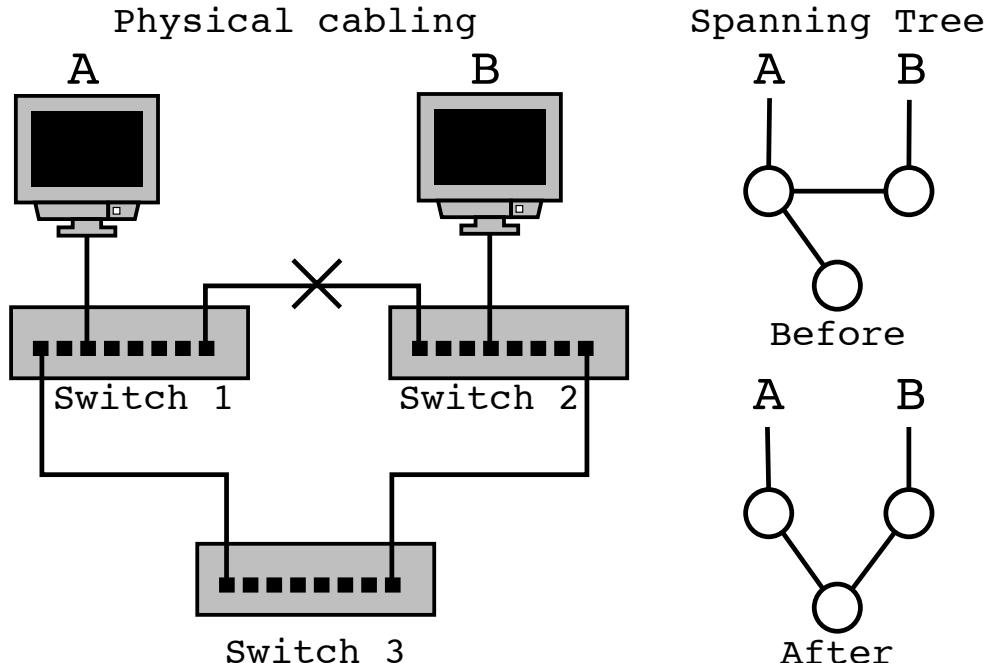
- Node 1: University of California Los Angeles
- Node 2: Stanford Research Institute
- Node 3: University of California Santa Barbara
- Node 4: University of Utah

A switch



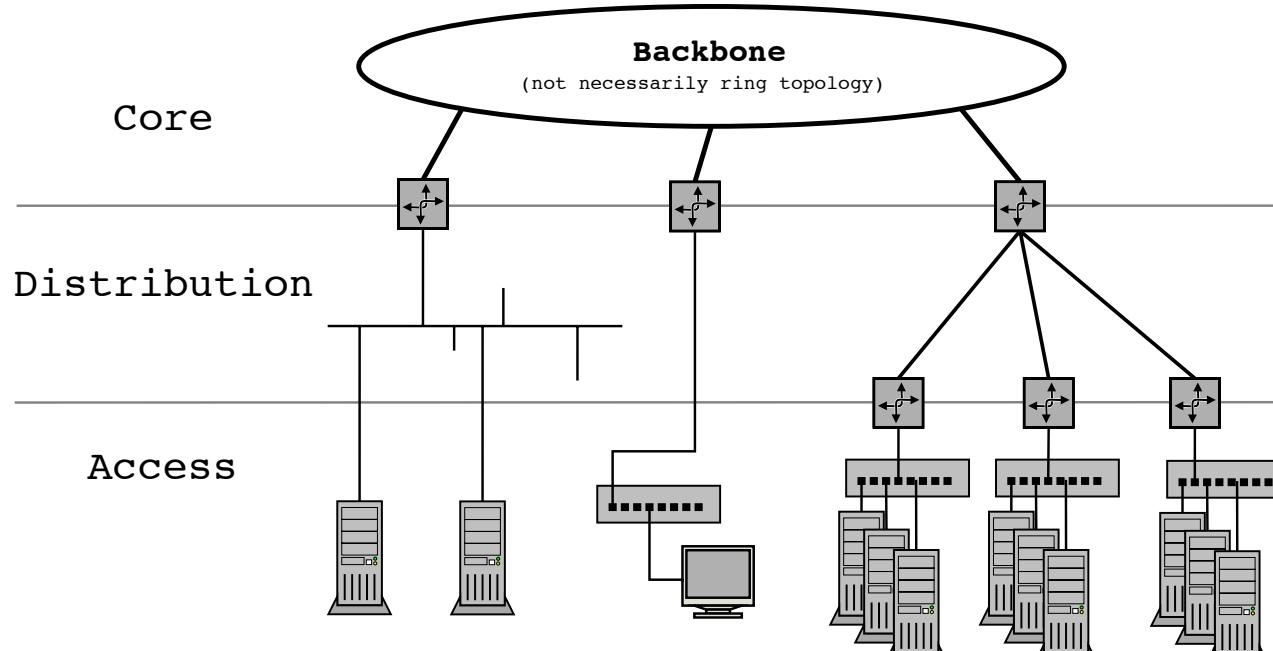
Today we use switches, Don't buy a hub, not even for experimenting or sniffing
A switch can receive and send data on multiple ports at the same time
Performance only limited by the backplane and switching chips
Can also often route with the same speed
Always buy managed switches with SNMP and IEEE 802.1q, even for home

Topologier og Spanning Tree Protocol



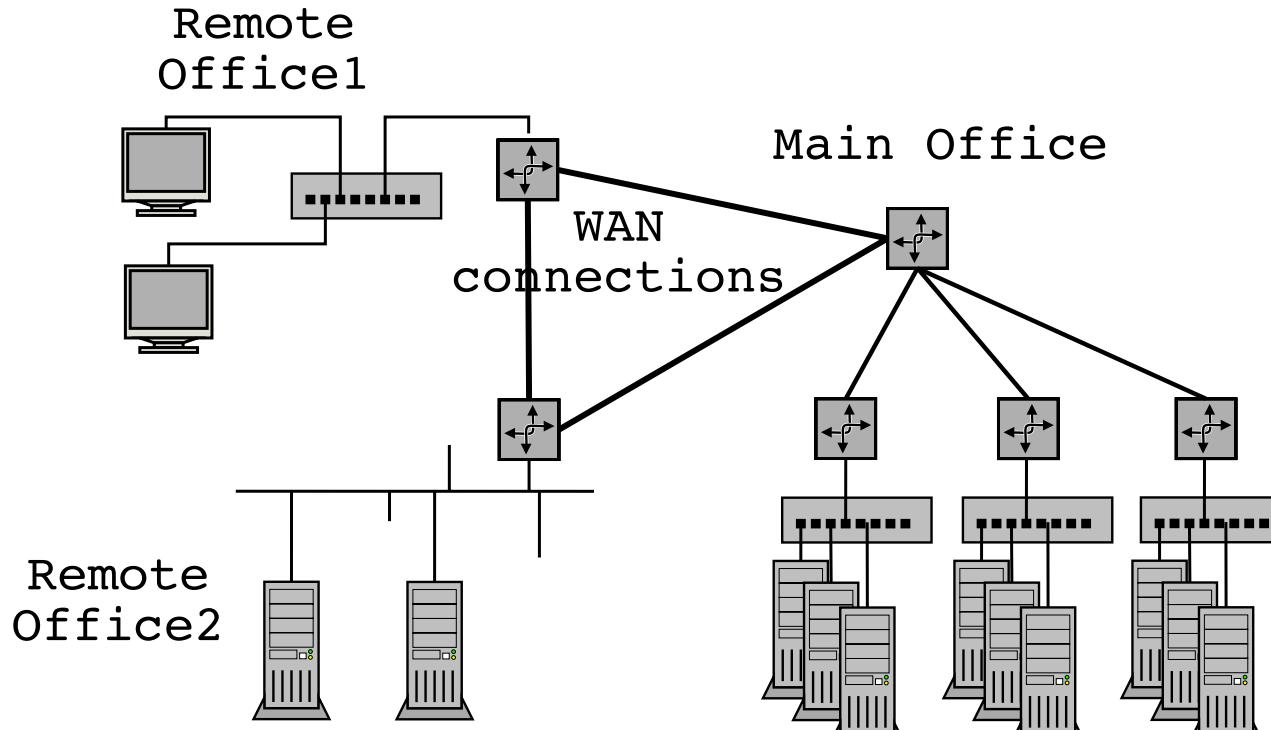
See also Radia Perlman, *Interconnections: Bridges, Routers, Switches, and Internetworking Protocols*

Core, Distribution og Access net



Not always this way - but often

Bridges and routers

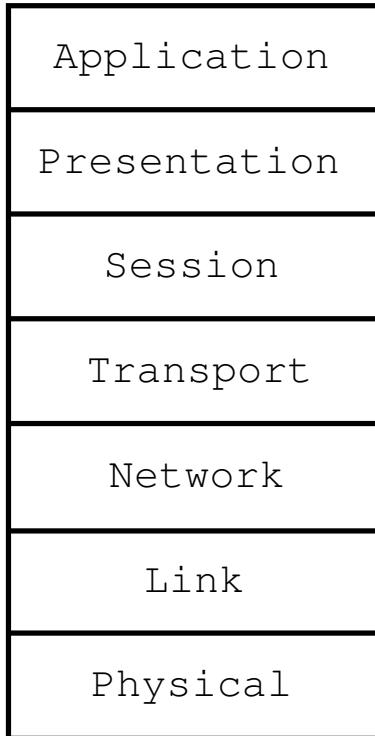


Often you inherit a mess

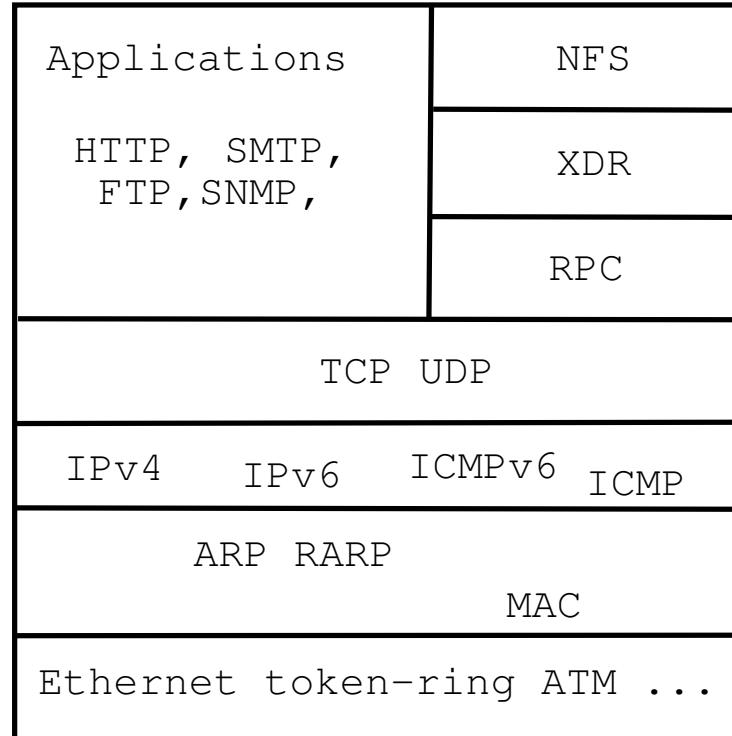
OSI and Internet models



OSI Reference Model



Internet protocol suite



MAC address



00-03-93	(hex)	Apple Computer, Inc.
000393	(base 16)	Apple Computer, Inc.
		20650 Valley Green Dr.
		Cupertino CA 95014
		UNITED STATES

Network technologies use a layer 2 hardware address

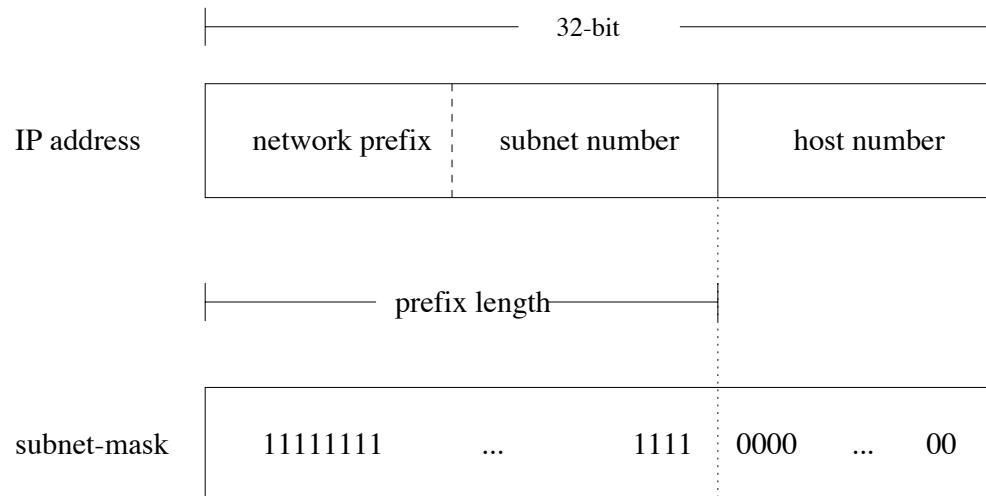
Typically using 48-bit MAC addresses known from Ethernet MAC-48/EUI-48

First half is assigned to companies – Organizationally Unique Identifier (OUI)

Using the OUI you can see which producer and roughly when a network chip was produced

<http://standards.ieee.org/regauth/oui/index.shtml>

Common Address Space



- Internet is defined by the address space, one
- Based on 32-bit addresses, example dotted decimal format 10.0.0.1

CIDR Classless Inter-Domain Routing



Classful routing		Classless routing CIDR	
4 class C networks	Inherent subnet mask	Supernet	Subnet mask
192.0.8.0	255.255.255.0	192.0.8.0	255.255.252.0
192.0.9.0	255.255.255.0		252d=11111100b
192.0.10.0	255.255.255.0		
192.0.11.0	255.255.255.0	Base network/prefix	
			192.10.8.0/22

- Subnet mask originally inferred by the class
- Started to allocate multiple C-class networks - save remaining B-class
Resulted in routing table explosion
- A subnet mask today is a row of 1-bit
- 10.0.0.0/24 means the network 10.0.0.0 with subnet mask 255.255.255.0
- Supernet, supernetting

Preparing an IPv6 Addressing Plan



Preparing an IPv6 Addressing Plan Manual

December 2010: Original text
March 2011: Translation provided by RIPE NCC



http://www.ripe.net/training/material/IPv6-for-LIRs-Training-Course/IPv6_addr_plan4.pdf

See also <https://blog.apnic.net/2019/08/22/how-to-ipv6-address-planning/>

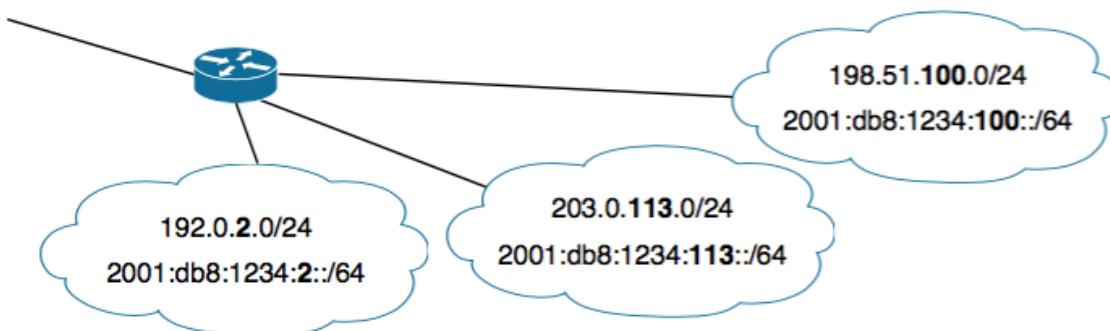


Example address plan input

3.1. Direct Link Between IPv4 and IPv6 Subnets

If the existing IPv4 networks use only /24 subnets (for example, from 203.0.113.0 to 203.0.113.255), a direct link can be established between IPv4 addresses and the new IPv6 addresses. In this case, you can include the penultimate number of the IPv4 address (113 in 203.0.113.0/24, for example) in the IPv6 subnet. The IPv6 address will then be 2001:db8:1234:113::/64.

Such an IPv4-to-IPv6 transition could appear as follows:



Easy and coupled with VLAN IDs it will work ☺

IP Address Management IPAM



NIPAP

127.0.0.1:5000/prefix/list#query_string=test&search_opt_parent=undefined&search_opt_child=undefined&explicit=false

VRF + VRFs prefixes pools Log out

test

Query took 0.64 seconds.

Search interpretation test: text matching 'test'.

Add prefix

VRF	Prefix	Order	FQDN	Description
No VRF	+ 1.0.0.0/8			
	+ 1.0.0.0/16			klonk
	1.0.1.0/24			test
	- 1.0.5.0/24			bla bla bla4
	1.0.5.1/24			test host 1
	1.0.5.2/24			test host 2
	1.0.5.3/24			test host 3
	1.0.5.4/24			test host 4
	1.0.5.5/24			test host 5
	1.0.5.6/24			test host 6
1.0.5.7/24			test host 7	
- 1.3.0.0/16			bla bla	
1.3.0.0/24			test	
1.3.3.0/24			blahona	
2.0.1.0/24			test	
2.0.5.0/24			test	
2.0.6.0/24			test	
2.0.7.0/24			test	
2.0.8.0/24			test	

http://127.0.0.1:5000/prefix/list#query_string=test&search_opt_parent=undefined&search_opt_child=undefined&explicit=false

- Recommend Nipap <http://spritelink.github.io/NIPAP/>

NIPAP example, adding firewall cluster interfaces - multiple VLANs

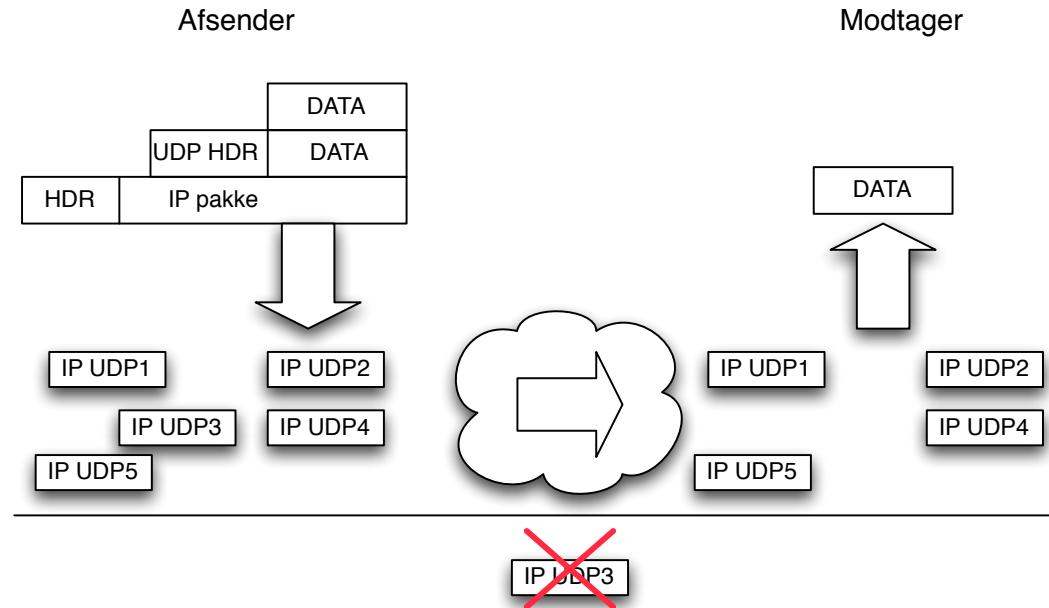


```
#!/bin/sh
# Add sites to NIPAP, dont waste time doing it with mouse
addsite () {
    SITE=$1
    SITEID=$2
    for VLAN in 100 200 300 400 500
    do
        nipap address add family ipv4 prefix 172.$SITEID.$VLAN.1 node $SITE-fw.example.net
        nipap address add family ipv4 prefix 172.$SITEID.$VLAN.2 node $SITE-fw-01.example.net
        nipap address add family ipv4 prefix 172.$SITEID.$VLAN.3 node $SITE-fw-02.example.net
    done
}

addsite dk-odense 1231
addsite dk-svendborg 1232
```

- Automating saves time, and less errors from manual input!

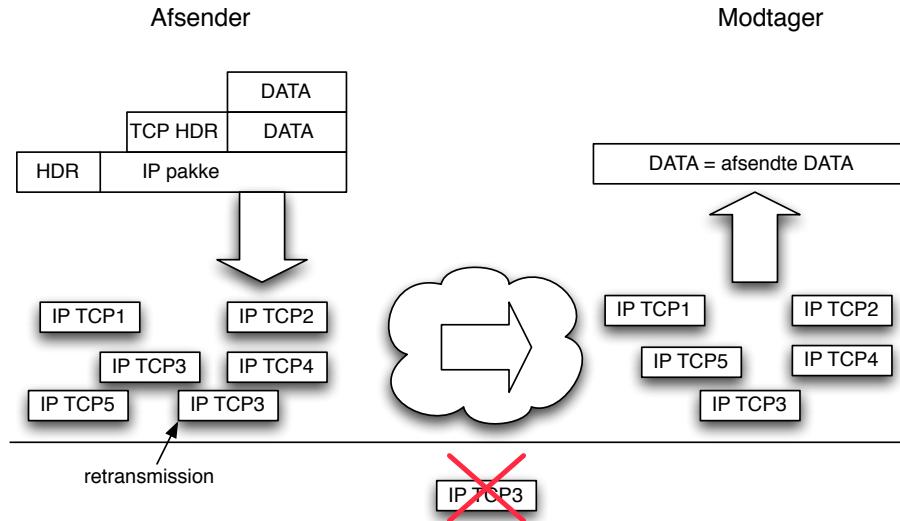
UDP User Datagram Protocol



Connection-less RFC-768, *connection-less*

Used for Domain Name Service (DNS)

TCP Transmission Control Protocol



Connection oriented RFC-791 September 1981, *connection-oriented*

Either data delivered in correct order, no data missing, checksum or an error is reported

Used for HTTP and others

Well-known port numbers



IANA maintains a list of magical numbers in TCP/IP
Lists of protocol numbers, port numbers etc.

A few notable examples:

- Port 25/tcp Simple Mail Transfer Protocol (SMTP)
- Port 53/udp and 53/tcp Domain Name System (DNS)
- Port 80/tcp Hyper Text Transfer Protocol (HTTP)
- Port 443/tcp HTTP over TLS/SSL (HTTPS)

Source: <http://www.iana.org>

ICMP Internet Control Message Protocol



Control protocol, error messages

Common messages

- ICMP ECHO, anyone there?
- Host unreachable
- Port unreachable

signaling

Defined in RFC-792

Don't block all ICMP – wrong!



ICMP messages to allow – Similar for IPv6

Type

- 0 = net unreachable;
- 1 = host unreachable;
- 2 = protocol unreachable;
- 3 = port unreachable;
- 4 = fragmentation needed and DF set;
- 5 = source route failed.

If you remove all ICMP you will have time outs instead

Allow these ICMP types:

- 3 Destination Unreachable
- 4 Source Quench Message
- 11 Time Exceeded
- 12 Parameter Problem Message

IPv6 neighbor discovery protocol (NDP)



OSI	IPv4	IPv6
Network	IP / ICMP	IPv6 / ICMPv6
Link	ARP	
Physical	Physical	Physical

Address Resolution Protocol (ARP) is gone from IPv6

NDP replaces and expands, command to use arp -an replaced by ndp -an

ARP vs NDP



```
hlk@bigfoot:basic-ipv6-new$ arp -an
? (10.0.42.1) at 0:0:24:c8:b2:4c on en1 [ethernet]
? (10.0.42.2) at 0:c0:b7:6c:19:b on en1 [ethernet]
hlk@bigfoot:basic-ipv6-new$ ndp -an
Neighbor          Linklayer Address  Netif Expire      St Flgs Prbs
::1                  (incomplete)        lo0 permanent R
2001:16d8:fffd2:cf0f:21c:b3ff:fec4:e1b6 0:1c:b3:c4:e1:b6 en1 permanent R
fe80::1%lo0          (incomplete)        lo0 permanent R
fe80::200:24ff:fec8:b24c%en1 0:0:24:c8:b2:4c       en1 8h54m51s S R
fe80::21c:b3ff:fec4:e1b6%en1 0:1c:b3:c4:e1:b6       en1 permanent R
```

Basic test tools TCP - Ping and Traceroute



We should all know

- ping – like sending a radar ping, anything there
- traceroute (windows tracert) – find the route packets traverse

and add these!

- Wireshark – like sending a radar ping, anything there
- Nmap and Nping – port scan and advanced ping program!

Ping



```
$ ping 192.168.1.1
PING 192.168.1.1 (192.168.1.1): 56 data bytes
64 bytes from 192.168.1.1: icmp_seq=0 ttl=150 time=8.849 ms
64 bytes from 192.168.1.1: icmp_seq=1 ttl=150 time=0.588 ms
64 bytes from 192.168.1.1: icmp_seq=2 ttl=150 time=0.553 ms
```

ICMP – Internet Control Message Protocol

ECHO – only ICMP message that generates another

ICMP ECHO request sent, and ICMP ECHO reply expected

Same with IPv6, ping6

traceroute



```
hkj@bob:~$ traceroute www.kramse.dk
traceroute to www.kramse.dk (185.129.60.130), 30 hops max, 60 byte packets
 1  10.0.42.1 (10.0.42.1)  0.365 ms  0.277 ms  0.239 ms
 2  79.142.xxx.xxx (79.142.xxx.xxx)  5.174 ms  4.979 ms  5.113 ms
 3  bgp2-dix.prod.bolignet.dk (79.142.224.2)  5.538 ms  5.057 ms  5.483 ms
 4  217.74.215.57 (217.74.215.57)  5.990 ms  5.962 ms  5.932 ms
...
 8  185.150.199.178 (185.150.199.178)  7.684 ms  7.647 ms  4.627 ms
 9  * * * // firewall here!
```

Works using the Time to live (TTL) counter

Sending with $TTL = 1$ returns ICMP from first host/router

Default sends UDP on Unix, and ICMP on Windows

Kali has programs that can emulate, or send using any protocol

Wireshark - graphical network sniffer



We're having a conference! You're invited!

Download
Get Started Now

Learn
Knowledge is Power

Enhance
With Riverbed Technology

News And Events

Join us at SHARKFEST '15!
SHARKFEST '15 will be held from June 22 – 25 at the Computer History Museum in Mountain View, CA.
[Learn More ▶](#)

Troubleshooting with Wireshark
By Laura Chappell
Foreword by Gerald Combs
Edited by Jim Aragon

This book focuses on the tips and techniques used to identify

Wireshark Blog

Cool New Stuff
Dec 17 | By Evan Huus

Wireshark 1.12 Officially Released!
Jul 31 | By Evan Huus

To Infinity and Beyond! Capturing Forever with Tshark
Jul 8 | By Evan Huus
[More Blog Entries ▶](#)

Enhance Wireshark

Riverbed is Wireshark's primary sponsor and provides our funding. They also make great products.

802.11 Packet Capture

- WLAN packet capture and transmission
- Full 802.11 a/b/g/n support
- View management, control and data frames
- Multi-channel aggregation (with multiple adapters)

[Learn More ▶](#)

[Buy Now ▶](#)

<http://www.wireshark.org>

Using Wireshark



```
http-example.cap
```

Apply a display filter: <ctrl>F

No.	A	Time	Source	Destination	Protocol	Info
1	0.000000	172.24.65.102	91.182.91.18	TCP	58816 - http	[SYN] Seq=0 Win=65535 Len=0 MSS=1460 WS=16 TSval=745562412 TSecr=0 SACK_PERM=0
2	0.000170	172.24.65.102	91.182.91.18	TCP	58817 - http	[SYN] Seq=0 Win=65535 Len=0 MSS=1460 WS=16 TSval=745562412 TSecr=0 SACK_PERM=0
3	0.127053	91.182.91.18	172.24.65.102	TCP	58816 - http	[SYN, ACK] Seq=0 Ack=1 Win=16384 Len=0 MSS=1460 SACK_PERM=1 WS=8 TSval=1855239975
4	0.127167	91.182.91.18	172.24.65.102	TCP	58817 - http	[SYN, ACK] Seq=0 Ack=1 Win=16384 Len=0 MSS=1460 SACK_PERM=1 WS=8 TSval=2512433851
5	0.127181	172.24.65.102	91.182.91.18	TCP	58816 - http	[ACK] Seq=1 Ack=1 Win=131760 Len=0 TSval=45562538 TSecr=1855239975
6	0.127226	172.24.65.102	91.182.91.18	TCP	58817 - http	[ACK] Seq=1 Ack=1 Win=131760 Len=0 TSval=45562538 TSecr=2512433851
7	0.127363	172.24.65.102	91.182.91.18	HTTP	GET / HTTP/1.1	
8	0.141320	91.182.91.18	172.24.65.102	HTTP	HTTP/1.1 304 Not Modified	
9	0.141421	172.24.65.102	91.182.91.18	TCP	58816 - http	[ACK] Seq=583 Ack=190 Win=131568 Len=0 TSval=745562551 TSecr=1855239975

Frame 7: 568 bytes on wire (4544 bits), 568 bytes captured (4544 bits)
Ethernet II, Src: Apple_E6:87:5e (7c:01:c3:6c:87:5e), Dst: Cisco_32:09:38 (44:2b:03:32:09:38)
Internet Protocol Version 4, Src: 172.24.65.102 (172.24.65.102), Dst: 91.182.91.18 (91.182.91.18)
Transmission Control Protocol, Src Port: 58816 (58816), Dst Port: http (80), Seq: 1, Ack: 1, Len: 502
HyperText Transfer Protocol
 > GET / HTTP/1.1\r\n Host: 91.182.91.18\r\n Connection: keep-alive\r\n Cache-Control: max-age=0\r\n Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/webp,*/*;q=0.8\r\n User-Agent: Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_2) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/33.0.1750.146 Safari/537.36\r\n Accept-Encoding: gzip,deflate,sdch\r\n Accept-Language: en-US,en;q=0.8,cs;q=0.6,da;q=0.4\r\n If-None-Match: 7053a63e1516a827a95db3107524a6e0a3\r\n If-Modified-Since: Tue, 17 Nov 2009 11:22:22 GMT\r\n Vary: *\r\n [HTTP request URL: http://91.182.91.18/] [HTTP request 1/1] [Response in frame: 8]

0000 44 2b 03 32 09 38 7c d1 c3 6c 07 5e 00 00 45 00 D+ 2.0 | ÄL...E.
0001 02 2b 0e d7 00 00 00 06 5f ff ac 18 41 66 5d 00 .*..x@. ØÙ~.T
0002 52 12 c5 e0 00 00 08 0e 0c c7 03 14 0c 19 80 18 l.ÄÄ,P.É.Ç.....
0038 2b 2b 0f c0 00 00 00 01 00 00 00 00 00 00 00 00 +.Ä... ...,.pañ.
0048 b7 27 47 54 28 2f 28 48 54 58 0f 2f 31 2e 31 .GET / HTTP/1.1
0069 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .Host: 91.182.9
0070 3b 2b 6b 65 65 70 2d 61 6e 69 76 65 0d 0a 43 61 .1.Ä...Co...
0088 63 68 65 24 43 66 74 72 6f 6c 3a 20 6d 61 78 che-Cont ro: max
0098 2d 61 67 65 3d 08 0d 41 63 63 65 70 74 3a 20 -age=0. Accept:
0080 75 65 78 74 2f 68 74 6d 6c 2d 61 70 78 6c 69 63 text/html, applic
0061 64 74 69 67 6c 2f 78 68 74 6d 6c 2b 78 6d 6c 2c application/xm
0082 61 78 70 6e 69 63 61 74 69 6f 6e 2f 78 6d 6c 3b application/xm;

Capture - Options, select a network interface

Detailed view of network traffic with Wireshark



```
▼ Secure Sockets Layer
  ▼ TLSv1.2 Record Layer: Handshake Protocol: Client Hello
    Content Type: Handshake (22)
    Version: TLS 1.0 (0x0301)
    Length: 198
  ▼ Handshake Protocol: Client Hello
    Handshake Type: Client Hello (1)
    Length: 194
    Version: TLS 1.2 (0x0303)
    ▶ Random
      Session ID Length: 0
      Cipher Suites Length: 32
    ▶ Cipher Suites (16 suites)
    ▶ Compression Methods Length: 1
    ▶ Compression Methods (1 method)
    ▶ Extensions Length: 121
    ▶ Extension: Unknown 56026
    ▶ Extension: renegotiation_info
    ▼ Extension: server_name
      Type: server_name (0x0000)
      Length: 16
      ▶ Server Name Indication extension
        Server Name list length: 14
        Server Name Type: host_name (0)
        Server Name length: 11
        Server Name: twitter.com
      ▶ Extension: Extended Master Secret
0050 a4 1d 52 8f 2c 18 99 91 54 68 0a 77 0d 95 73 64 ..R,... Th.w..sd
0060 7d 00 00 20 5a 5a c0 2b c0 2f c0 2c c0 30 cc a9 }.. ZZ.+ ./.,.0..
0070 cc a8 cc 14 cc 13 c0 13 c0 14 00 9c 00 9d 00 2f ..... ....../
0080 00 35 00 0a 01 00 00 79 da da 00 00 ff 01 00 01 .5.....y .....
0090 00 00 00 10 00 0e 00 00 0b 74 77 69 74 74 65 ..... ..twitte
00a0 72 2e 63 6f 6d 00 17 00 00 00 23 00 00 00 0d 00 r.com... ..#.....
00b0 14 00 12 04 03 08 04 04 01 05 03 08 05 05 01 08 ..... .....
```

Notice also the filtering possibilities, capture and view

Remote network debugging



- TShark and Tcpdump, I often use:
`tcpdump -nei eth0`
`tshark -z expert -r download-slow.pcapng`
- Remote packet dumps, `tcpdump -i eth0 -w packets.pcap`
- Story: tcpdump was originally written in 1988 by Van Jacobson, Sally Floyd, Vern Paxson and Steven McCanne who were, at the time, working in the Lawrence Berkeley Laboratory Network Research Group
<https://en.wikipedia.org/wiki/Tcpdump>

Great network security comes from knowing networks!

Note About Hardware IPv4 checksum offloading



IPv4 checksum must be calculated for every packet received

IPv4 checksum must be calculated for every packet sent

Usually on a router the Time To Live is decremented, to need re-calculation

Let an ASIC chip on the network card do the work!

Most server network chips today support this and more

Benefit for performance, but beware when using security tools

If every packet in wireshark has wrong checksum, its the network card doing it

Can be turned off, when doing security work

Nping



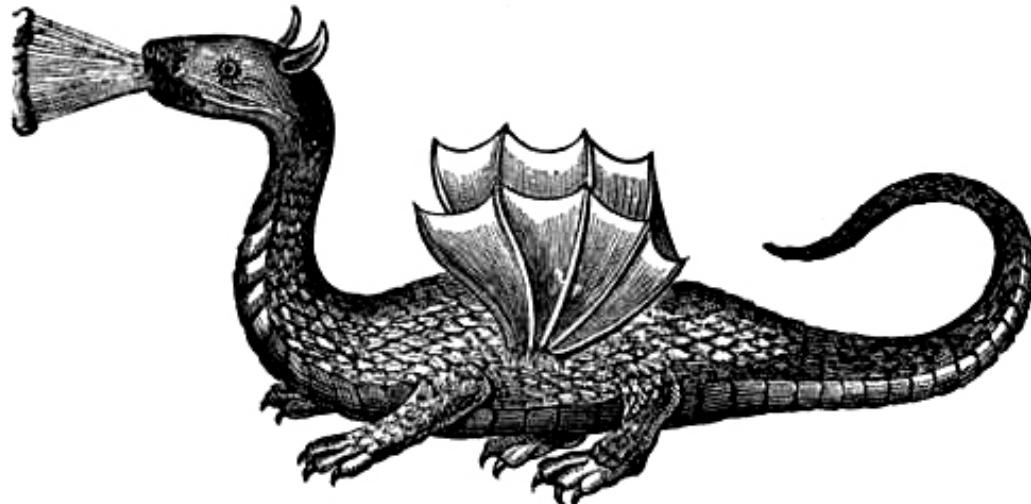
```
root@KaliVM:~# nping --tcp -p 80 www.zecurity.com
```

```
Starting Nping 0.7.70 ( https://nmap.org/nping ) at 2018-09-07 19:06 CEST
SENT (0.0300s) TCP 10.137.0.24:3805 > 185.129.60.130:80 S ttl=64 id=18933 iplen=40 seq=2984847972 win=1480
RCVD (0.0353s) TCP 185.129.60.130:80 > 10.137.0.24:3805 SA ttl=56 id=49674 iplen=44 seq=3654597698 win=16384 <mss 1460>
SENT (1.0305s) TCP 10.137.0.24:3805 > 185.129.60.130:80 S ttl=64 id=18933 iplen=40 seq=2984847972 win=1480
RCVD (1.0391s) TCP 185.129.60.130:80 > 10.137.0.24:3805 SA ttl=56 id=50237 iplen=44 seq=2347926491 win=16384 <mss 1460>
SENT (2.0325s) TCP 10.137.0.24:3805 > 185.129.60.130:80 S ttl=64 id=18933 iplen=40 seq=2984847972 win=1480
RCVD (2.0724s) TCP 185.129.60.130:80 > 10.137.0.24:3805 SA ttl=56 id=9842 iplen=44 seq=2355974413 win=16384 <mss 1460>
SENT (3.0340s) TCP 10.137.0.24:3805 > 185.129.60.130:80 S ttl=64 id=18933 iplen=40 seq=2984847972 win=1480
RCVD (3.0387s) TCP 185.129.60.130:80 > 10.137.0.24:3805 SA ttl=56 id=1836 iplen=44 seq=3230085295 win=16384 <mss 1460>
SENT (4.0362s) TCP 10.137.0.24:3805 > 185.129.60.130:80 S ttl=64 id=18933 iplen=40 seq=2984847972 win=1480
RCVD (4.0549s) TCP 185.129.60.130:80 > 10.137.0.24:3805 SA ttl=56 id=62226 iplen=44 seq=3033492220 win=16384 <mss 1460>
```

```
Max rtt: 40.044ms | Min rtt: 4.677ms | Avg rtt: 15.398ms
Raw packets sent: 5 (200B) | Rcvd: 5 (220B) | Lost: 0 (0.00%)
Nping done: 1 IP address pinged in 4.07 seconds
```

- The Nmap portscanner includes a nice ping utility, that allows us to *ping* with TCP

Challenges in network management



Internet here be dragons

We will jump directly into some solutions

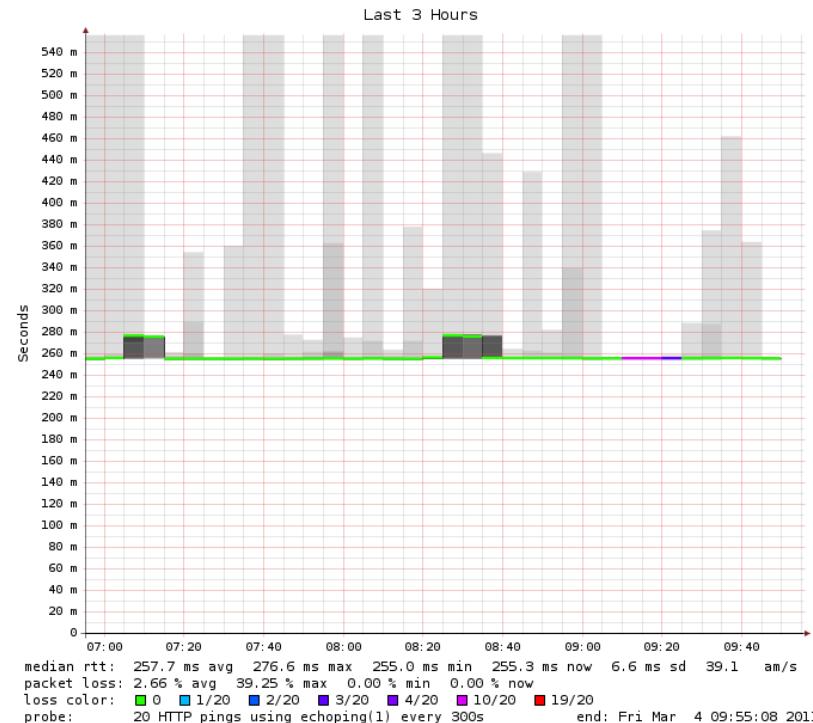
The basic tools for monitoring network



Moving from manual checking to automated

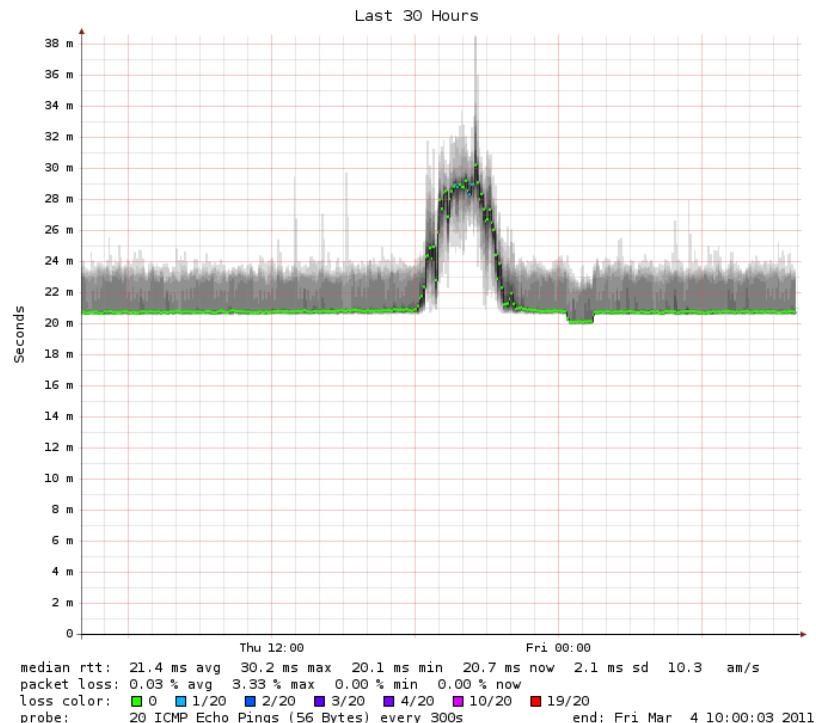
- Wireshark is advanced manual tool, try right-clicking different places
- How do we continuously monitor the network?

Smokeping packet loss



Old skool, but very usefull <https://oss.oetiker.ch/smokeping/>

Smoking latency changed



NTP Network Time Protocol



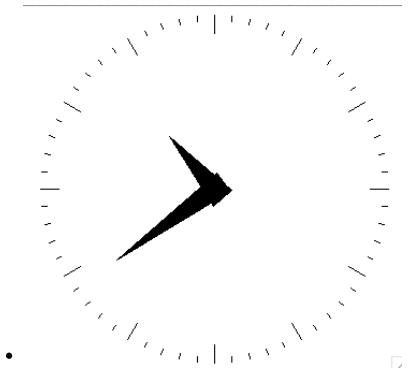
Vigtigt at netværksenheder bruger korrekt tid, sikkerhed og drift

Server NTP foregår typisk i /etc/ntp.conf eller /etc/ntpd.conf
det vigtigste er navnet på den/de servere man vil bruge som tidskilde

Brug enten en NTP server hos din udbyder eller en fra <http://www.pool.ntp.org/>

Eksempelvis:

```
server 0.dk.pool.ntp.org
server 0.europe.pool.ntp.org
server 3.europe.pool.ntp.org
```



SNMP and management



Often we see devices in the network configured using HTTP, Telnet eller SSH

- Where is the configuration stored?
- Is the device functioning?

Today we can use automation like:

- (old RANCID <http://www.shrubbery.net/rancid/>)
- Ansible <https://www.ansible.com/>
- Python with SSH libraries
- Oxidized is a network device configuration backup tool. RANCID replacement! <https://github.com/ytti/oxidized>

SNMP version 2 vs version 3



```
snmpwalk -v2c -c public 172.16.1.2 .1.3.6.1.4.1 | grep site
SNMPv2-SMI::enterprises.13742.6.3.2.2.1.13.1 = STRING: "site-pdu-rack-a-27"
SNMPv2-SMI::enterprises.13742.6.3.5.3.1.3.1.1 = STRING: "switch site-dist-sw-02"
SNMPv2-SMI::enterprises.13742.6.3.5.3.1.3.1.11 = STRING: "router site-mx-01"
```

Use Simple Network Management Protocol, but

- Typical values, packet count on ports, errors observed, speed, firmware, versions,
- SNMP versions 1 and 2c are insecure
- SNMP version 3 created to fix this
- Even better put SNMP and management in separate VLAN

Example for Juniper can be found at:

https://www.juniper.net/documentation/en_US/junos/topics/example/snmpv3-configuration-junos-nm.html

Config example: SNMP



```
snmp {  
    description "CORE-SW-02";  
    location "Teknikrum, Graested, Denmark";  
    contact "noc@zencurity.com";  
    community yourcommunitynotmine {  
        authorization read-only;  
        clients {  
            10.1.1.1/32;  
            10.1.2.2/32;  
        }  
    }  
}
```

SNMPv2 example

RANCID output



carrier-switches router config diffs — RANCID

From: rancid@[REDACTED]
Subject: carrier-switches router config diffs
Date: 19. jan 2011 02.09.28 CET
To: rancid-carrier-switches@[REDACTED]

Index: configs/c1-cph-01

```
-- configs/c1-cph-01 (revision 1457)
@@ -210,7 +210,7 @@
exit
!
interface ethernet 1/g45
- description 'SRX240'
+ description 'SRX-CPH-02 ge-0/0/0'
switchport mode general
switchport general allowed vlan add 95,3000-3005 tagged
exit
```

Step 1: configure devices properly



You should always configure your devices properly

Use Secure Shell (SSH)

Configure NTP

Turn on SNMP, probably SNMPv3

Turn on LLDP Link Layer Discovery Protocol – vendor-neutral

http://en.wikipedia.org/wiki/Link_Layer_Discovery_Protocol

Configure centralized syslog

And updated firmware, HTTPS and SSH only etc. the usual stuff

Then use Oxidized and LibreNMS

Oxidized



```
gem install oxidized  
gem install oxidized-script oxidized-web
```

- Oxidized is written in Ruby as a replacement for RANCID (Perl)
- Very easy to get running
- Fetches and imports configurations into Git repository, or others
- RESTful API and Web Interface, and command line tools
- <https://github.com/ytti/oxidized>

Oxidized configuration



```
~/.config/oxidized/router.db  
router01.example.com:ios  
switch01.example.com:procurve  
router02.example.com:ios
```

Easy to configure new device types, added Clavister firewalls once



Config example: Dell switch LLDP

```
interface ethernet 1/xg17
mtu 9216
lldp transmit-tlv port-desc sys-name sys-desc sys-cap
lldp transmit-mgmt
exit
```



LLDP trick using tcpdump

```
[hlk@ljh ~]$ sudo tcpdump -i eth0 ether proto 0x88cc
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on eth0, link-type EN10MB (Ethernet), capture size 96 bytes
.... wait for it ....
11:03:55.395064 00:1c:23:80:49:ff (oui Unknown) > 01:80:c2:00:00:0e (oui Unknown),
ethertype Unknown (0x88cc), length 60:
0x0000: 0207 0400 1c23 8049 fd04 0705 312f 302f  ....#.I....1/0/
0x0010: 3331 0602 0078 0000 0000 0000 0000 0000  31...x.....
0x0020: 0000 0000 0000 0000 0000 0000 0000      .....
1 packets captured
3 packets received by filter
0 packets dropped by kernel
```

I know **for sure** that this server is in Unit 1 port 31!

LibreNMS a fully featured network monitoring system



- Homepage: <https://www.librenms.org/>
- **Suggested method:** <https://docs.librenms.org/Installation/>
- How basic information about devices are presented, from devices when added - and nothing more.
See how to add a device and add your own. <https://docs.librenms.org/Support/Adding-a-Device/>
- How SNMP location is used to categorize devices and provide maps, see
<https://docs.librenms.org/Extensions/World-Map/>
- How protocols like LLDP allow LibreNMS to make maps, see
<https://docs.librenms.org/Extensions/Network-Map/>
- How port description can be used for describing ports
<https://docs.librenms.org/Extensions/Interface-Description-Parsing/>

Most of this happens with very little effort. Just configure devices consistently and they will be presented nicely.

LibreNMS Automatic discovery



LibreNMS

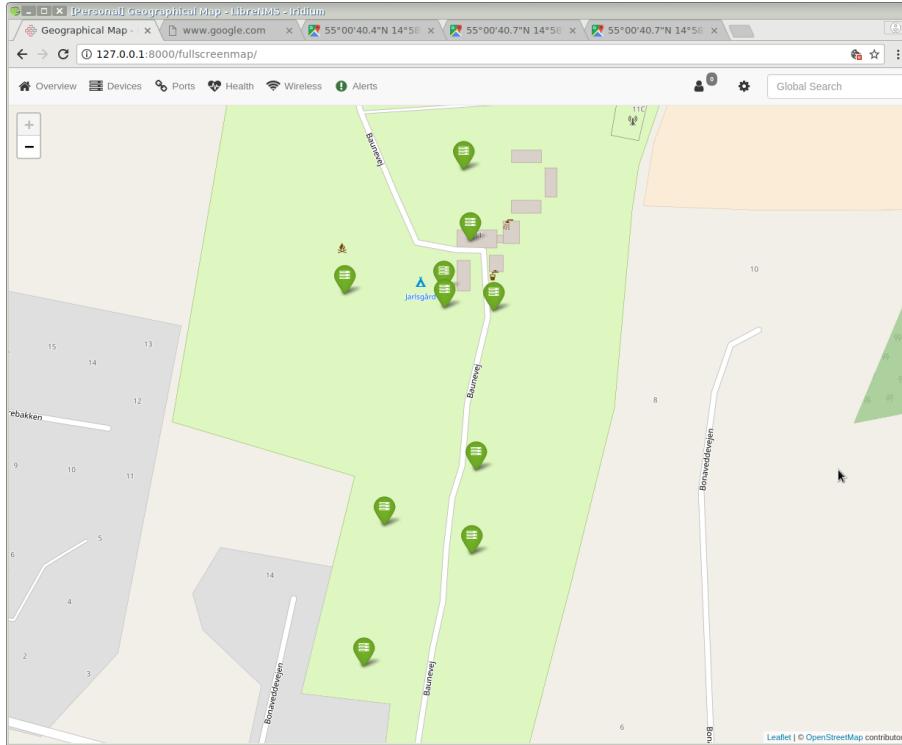
Overview Devices Ports Health Wireless Alerts

Lists: Basic | **Detail** Graphs: Bits | CPU | Load | Memory | Uptime | Storage | Disk I/O | Poller | Ping | Temperature

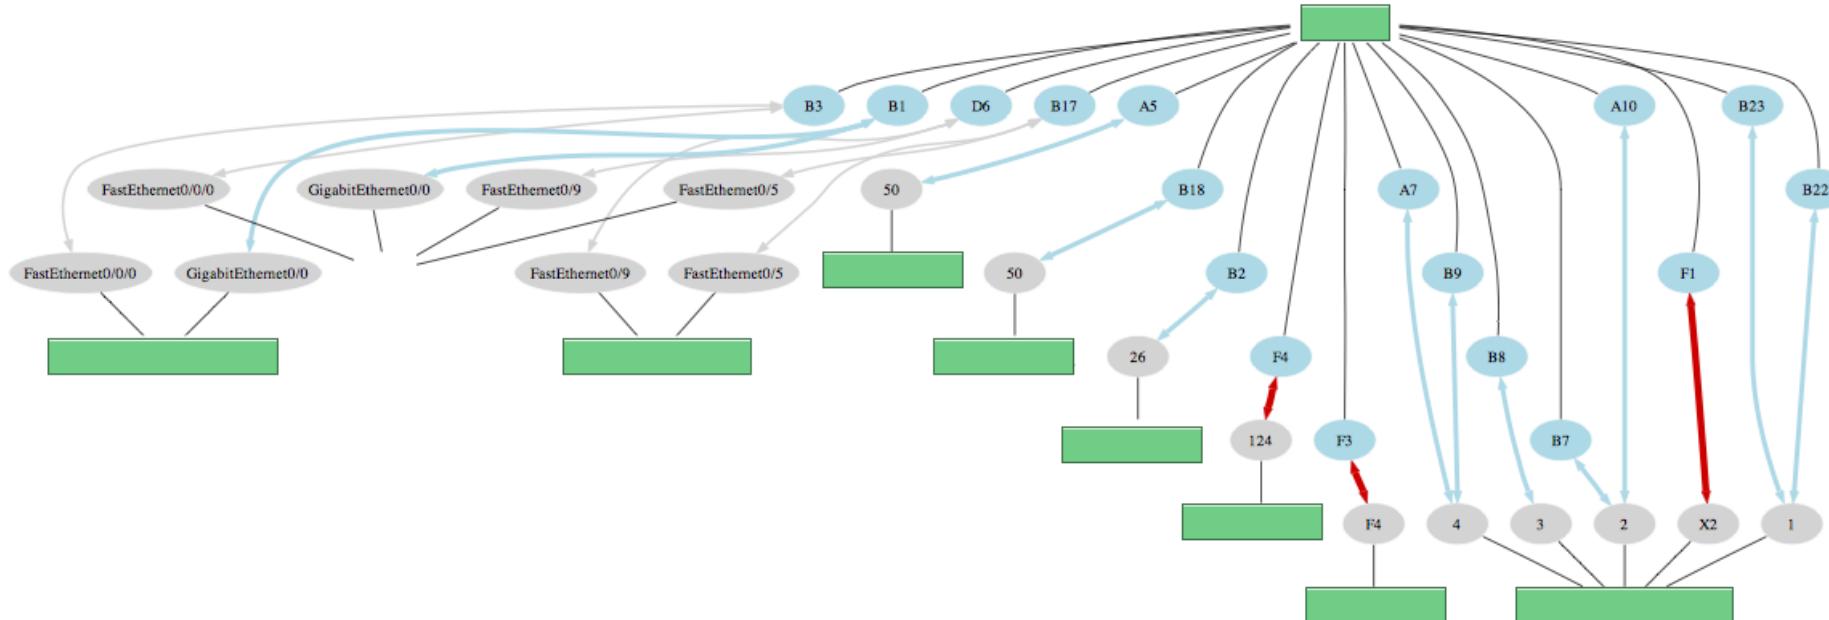
Vendor	Device	Metrics	Platform	Operating System
	192.168.0.254 zw-zd3k-001	7 2	zd3025	Ruckus Wireless 10.1.1.0 build 42 (DK)
	born-core-01	102 13	Juniper EX3300	Juniper JunOS 15.1R2.9
	noctent1 noc-tent	29 3	Brocade ICX 6430 24-port Switch	Brocade IronWare
	north1 north1	25		Foundry Networking
	south1 south1	25 4	snFWS624GSwitch	Brocade IronWare
	south2 south2	29 3	Brocade ICX 6430 24-port Switch	Brocade IronWare
	south3 south3	49		Foundry Networking
	southwest1 southwest1	49		Foundry Networking
	west1 west1	25 4	snFWS624GSwitch	Brocade IronWare
	west2 west2	25		Foundry Networking

Automatically discover your entire network using CDP, FDP, LLDP, OSPF, BGP, SNMP and ARP.

LibreNMS Geo Location



LLDP spaghetti?



LLDP is needed!

LibreNMS wireless clients



Centralized management SSH, Jump hosts



A jump server, jump host or jumpbox is a computer on a network used to access and manage devices in a separate security zone. The most common example is managing a host in a DMZ from trusted networks or computers.

https://en.wikipedia.org/wiki/Jump_server

Advantage, you can configure this server/system to only allow Key based logins – eliminate the possibility for brute-force attacks succeeding



OpenSSH client config with jump host

My recommended SSH client settings, put in \$HOME/.ssh/config:

```
Host *
  ServerAliveInterval=30
  ServerAliveCountMax=30
  NoHostAuthenticationForLocalhost yes
  HashKnownHosts yes
  UseRoaming no
```

```
Host jump-01
  Hostname 10.1.2.3
  Port 12345678
```

```
Host fw-site-01 10.1.2.5
  User hlk
  Port 34
  Hostname 10.1.2.5
  ProxyCommand ssh -q -a -x jump-01 -W %h:%p
```

I configure fw using both hostname and IP,
then I can use name, and any program using IP get this config too

Ansible



From my course materials:

Ansible is great for automating stuff, so by running the playbooks we can get a whole lot of programs installed, files modified - avoiding the Vi editor.

- Easy to read, even if you don't know much about YAML
- <https://www.ansible.com/> and [https://en.wikipedia.org/wiki/Ansible_\(software\)](https://en.wikipedia.org/wiki/Ansible_(software))
- Great documentation
https://docs.ansible.com/ansible/latest/collections/ansible/builtin/apt_module.html

Ansible Dependencies



- Ansible based on Python, only need Python installed
<https://www.python.org/>
- Often you use Secure Shell for connecting to servers
<https://www.openssh.com/>
- Easy to configure SSH keys, for secure connections

How Ansible Works: inventory files



List your hosts in one or multiple text files:

```
[all:vars]
ansible_ssh_port=34443

[office]
fw-01 ansible_ssh_host=192.168.1.1 ansible_ssh_port=22
ansible_python_interpreter=/usr/local/bin/python

[infrastructure]
smtp-01      ansible_ssh_host=192.0.2.10
ansible_python_interpreter=/usr/local/bin/python
vpnmon-01    ansible_ssh_host=10.50.60.18
```

- Inventory files specify the hosts we work with
- Linux and OpenBSD servers shown here
- Real inventory for a site with development and staging may be 500 lines
- office and infrastructure are group names

How Ansible Works: ad hoc commands



Using the inventory file you can run commands with Ansible:

```
ansible -m ping new-server  
ansible -a "date" new-server  
ansible -m shell -a "grep a /etc/something" new-server
```

- Running commands on multiple servers is easy now
- This alone has value, you can start
- Checking settings on servers
- Making small changes to servers

How Ansible Works: Playbooks



The benefit comes with tasks listed in playbooks – example playbook content, installing software using APT:

```
apt:  
  name: "{{ packages }}"  
vars:  
  packages:  
    - nmap  
    - curl  
    - iperf  
    ...
```

Running it:

```
cd kramse-labs/suricatazeek  
ansible-playbook -v 1-dependencies.yml 2-suricatazeek.yml 3-elasticstack.yml
```

"YAML (a recursive acronym for "YAML Ain't Markup Language") is a human-readable data-serialization language."
<https://en.wikipedia.org/wiki/YAML>

Python and YAML – Git



- We need to store configurations
- Run playbooks
- Problem: Remember what we did, when, how
- Solution: use git for the playbooks
- Not the only version control system, but my preferred one

How Ansible Works: typical execution



```
ansible-playbook -i hosts.cph1 -K infrastructure-firewalls.yml -t pf.conf --check --diff
```

```
ansible-playbook -i hosts.cph1 -K infrastructure-firewalls.yml -t pf.conf
```

```
ansible-playbook -i hosts.cph1 -K infrastructure-nagios.yml -t config-only
```

```
ansible-playbook -i smartboxes -K create-pf-conf.yml -l smartbox-xxx-01
```

- Pro tip: check before you push out changes to production networks ☺
- Check will see if something needs changing
- Diff will show the changes about to be made
- Having configuration in Git improves things a lot!

Get ready, Up and running with Ansible



Prerequisites for Ansible - you need a Linux machine:

- python language - Ansible uses this
- ssh keys - remote login without passwords
- Sudo - allow regular users to do superuser tasks
- Recommended tool: ssh-copy-id for getting your key on new server
- Recommended Change: sshd_config - no passwords allowed, no bruteforce
- Recommended to use: jump hosts/ProxyCommand in ssh_config
- Highly recommended: Git and/or github for version control

Official docs:

https://docs.ansible.com/ansible/intro_installation.html

Ansible configuration management



```
- apt: name= item state=latest
  with_items:
    - unzip
    - elasticsearch
    - logstash
    - redis-server
    - nginx
- lineinfile: "dest=/etc/elasticsearch/elasticsearch.yml state=present
  regexp='network.host: localhost' line='network.host: localhost'"
- name: Move elasticsearch data into /data
  command: creates=/data/elasticsearch mv /var/lib/elasticsearch /data/
- name: Make link to /data/elasticsearch
  file: state=link src=/data/elasticsearch path=/var/lib/elasticsearch
```

Example playbooks:

```
git clone https://github.com/kramse/ansible-workshop
```

Ansible and Junos – Juniper Networks devices



Juniper Networks provides support for using Ansible to manage devices running Junos OS. Starting in Ansible Release 2.1, Ansible natively includes a number of core modules that can be used to manage devices running Junos OS. In addition, Juniper Networks provides the Juniper.junos Ansible role, which is hosted on the Ansible Galaxy website and includes additional modules for Junos OS.

Source: https://www.juniper.net/documentation/en_US/junos-ansible/topics/task/installation/junos-ansible-server-installing.html

- Minimal configuration changes: `user@host# set netconf ssh`

How cool is that, centralized configuration management with source code!

Logging



Logging

- The world needs answers
- You can only present answers, if you got them
- Using logging you can investigate what happened

Netflow – Lets move to detailed monitoring and logging



Cisco standard NetFlow version 5 defines a flow as a unidirectional sequence of packets that all share the following 7 values:

- Ingress interface (SNMP ifIndex), IP protocol, Source IP address and Destination IP address
- Source port for UDP or TCP, 0 for other protocols, IP Type of Service
- Destination port for UDP or TCP, type and code for ICMP, or 0 for other protocols

Today we can use Netflow version 9 or IPFIX with more fields available, or sFlow, short for "sampled flow", is an industry standard for packet export at Layer 2 of the OSI model,
<https://en.wikipedia.org/wiki/SFlow>

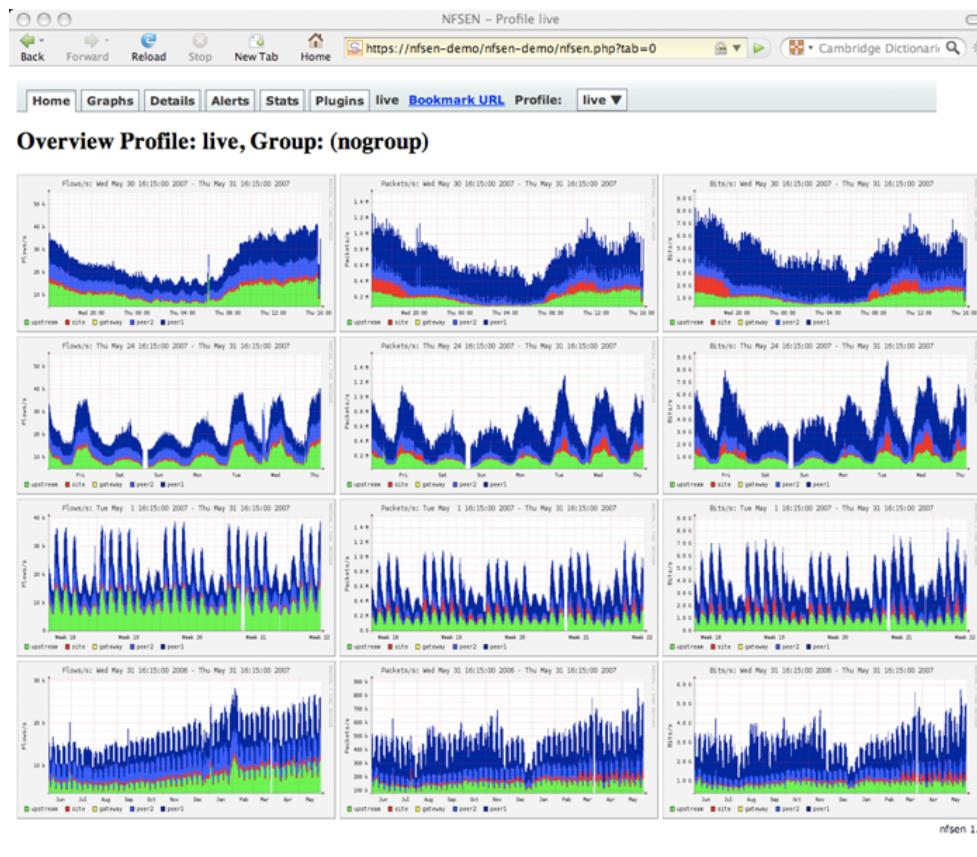
NFSen is an old but free application <http://nfsen.sourceforge.net/>

Source:

<https://en.wikipedia.org/wiki/NetFlow>

https://en.wikipedia.org/wiki/IP_Flow_Information_Export

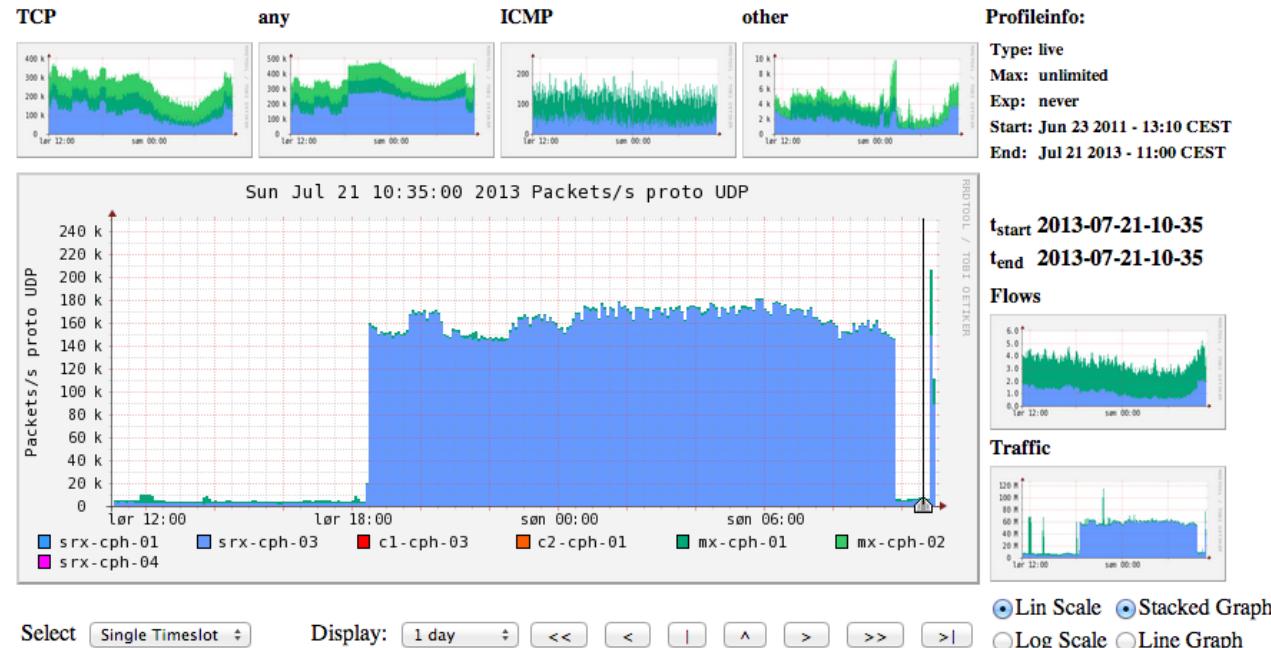
Netflow using NfSen



Netflow NFSen



Profile: live



An extra 100k packets per second from this netflow source (source is a router)

Netflow processing from the web interface



NFSEN - Profile live May 31 2007 - 04:40

Back Forward Reload Stop New Tab Home https://nfsen-demo/nfsen-demo/nfsen.php?processing Cambridge Dictionary

peer2 3.3 k/s 76.2 k/s 66.9 k/s 7.0 k/s 621.0 /s 1.7 k/s 484.6 Mb/s 459.9 Mb/s 12.5 Mb/s 437.3 kb/s 11.7 Mb/s
gateway 1.0 /s 651.0 /s 600.8 /s 46.6 /s 0 /s 3.7 /s 6.2 Mb/s 6.1 Mb/s 36.4 kb/s 0 b/s 4.4 kb/s
site 467.1 /s 8.9 k/s 6.1 k/s 2.0 k/s 181.7 /s 613.3 /s 38.8 Mb/s 28.3 Mb/s 7.4 Mb/s 104.0 kb/s 2.9 Mb/s
upstream 6.4 k/s 94.2 k/s 84.3 k/s 8.2 k/s 896.4 /s 766.7 /s 588.4 Mb/s 568.2 Mb/s 16.7 Mb/s 685.1 kb/s 2.8 Mb/s

All | None Display: Sum Rate

Netflow Processing

Source: peer1 Filter:

Options:

List Flows Stat TopN
Top: 10
Stat: Flow Records order by flows
proto
srcPort
dstPort
Aggregate
srcIP
dstIP
Limit: Packets > 0
Output: line / IPv6 long

Clear Form process

** nfdump -M /netflow0/nfsen-demo/profile-data/live/peer1:peer2:gateway:site:upstream -T -r 2007/05/31/04/nfcapd.200705310440
nfdump filter:
any
Aggregated flows 2797250
Top 10 flows ordered by flows:
Date flow start Duration Proto Src IP Addr:Port Dst IP Addr:Port Packets Bytes Flows
2007-05-31 04:39:54.045 299.034 UDP 116.147.95.88:1110 -> 188.142.64.162:27014 68 5508 68
2007-05-31 04:39:56.282 298.174 UDP 116.147.249.27:1478 -> 188.142.64.163:27014 67 5427 67
2007-05-31 04:39:57.530 298.206 UDP 117.196.44.62:1031 -> 188.142.64.166:27014 67 5427 67
2007-05-31 04:39:57.819 298.112 UDP 117.196.75.134:1146 -> 188.142.64.167:27014 67 5427 67
2007-05-31 04:39:53.187 297.216 UDP 61.191.235.132:4121 -> 60.9.138.37:4121 62 3720 62
2007-05-31 04:39:53.234 303.588 UDP 60.9.138.37:2121 -> 118.25.93.95:2121 61 3660 61
2007-05-31 04:39:58.921 298.977 UDP 60.9.138.37:2121 -> 121.135.4.186:2121 61 3660 61
2007-05-31 04:39:54.329 303.585 UDP 120.150.194.76:2121 -> 60.9.138.37:2121 61 3660 61
2007-05-31 04:39:53.916 300.734 UDP 60.9.138.37:2121 -> 125.167.25.128:2121 61 3660 61
2007-05-31 04:39:57.946 300.353 UDP 60.9.138.36:2121 -> 121.135.4.186:2121 61 3660 61

IP addresses anonymized
Summary: total flows: 4616424, total bytes: 156.6 G, total packets: 172.6 M, avg bps: 644.8 M, avg pps: 90946, avg bpp: 929
Time: 2007-05-31 04:44:55 - 2007-05-31 04:44:55
Total flows processed: 4616424, skipped: 0, Bytes read: 240064932
Sys: 6.184s flows/second: 746464.4 Wall: 6.185s flows/second: 746361.3

Bringing the power of the command line forward

ElastiFlow

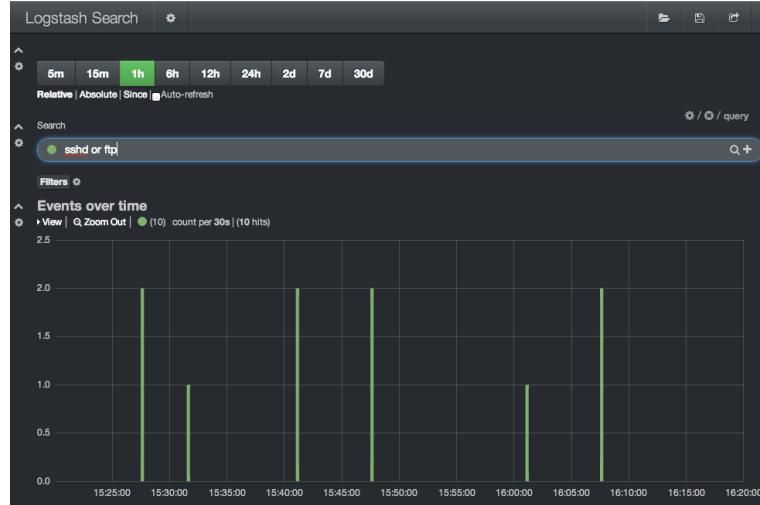


ElastiFlow™ provides network flow data collection and visualization using the Elastic Stack (Elasticsearch, Logstash and Kibana). It supports Netflow v5/v9, sFlow and IPFIX flow types (1.x versions support only Netflow v5/v9).

Source: Picture and text from <https://github.com/robcowart/elastiflow>

PS I havent tried it in real life, yet – but a friend has

View data efficiently

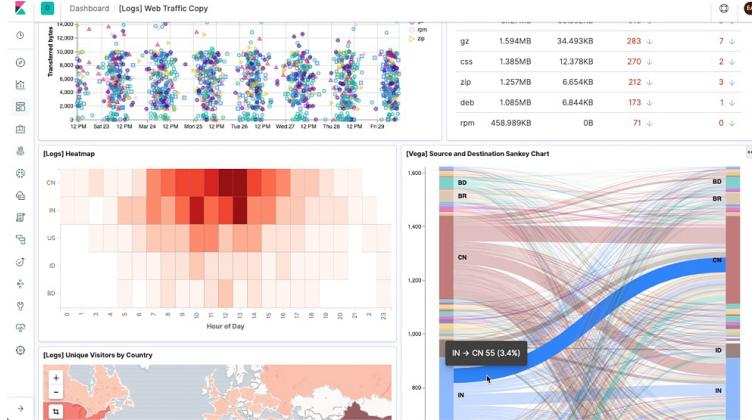


View data by digging into it easily - must be fast

Logstash and Kibana are just examples, but use indexing to make it fast!

Other popular examples include Graylog and Grafana

Big Data tools: Elasticsearch



Elasticsearch is an open source distributed, RESTful search and analytics engine capable of solving a growing number of use cases.

<https://www.elastic.co>

We are all Devops now, even security people!

Kibana

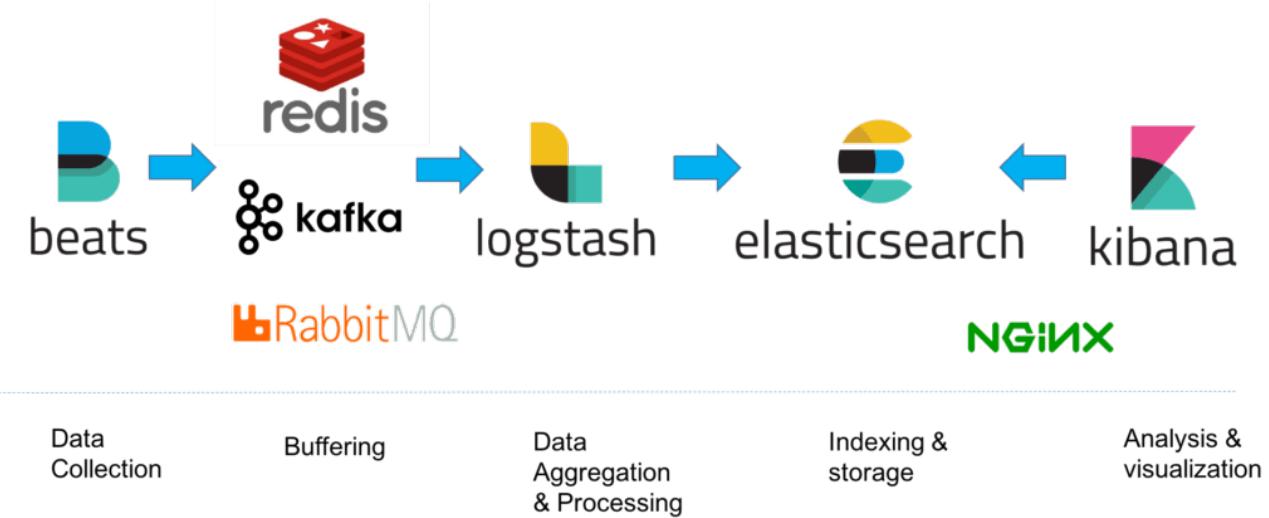


Highly recommended for a lot of data visualisation

Non-programmers can create, save, and share dashboards

Source: <https://www.elastic.co/products/kibana>

Architecture



- Real production environments often add some buffering in between
- Allows the ingestion to become more smooth, no lost messages

How to get started



How to get started searching for security events?

Collect basic data from your devices and networks

- Netflow data from routers
- Session data from firewalls
- Logging from applications: email, web, proxy systems

Centralize!

Process data

- Top 10: interesting due to high frequency, occurs often, brute-force attacks
- *ignore*
- Bottom 10: least-frequent messages are interesting

SIEM and logging systems



You could **buy a bunch of expensive gear**, point it all to a **log management** or a **security incident and event management (SIEM)** system, and let it automatically tell you what it knows. Some incident response teams may start this way, but unfortunately, many never evolve. **Working only with what the SIEM tells you**, versus what you have configured it to tell you **based on your contextual data**, will invariably fail. Truly demonstrating **value** from security monitoring and incident response **requires a major effort**. As with all projects, **planning** is the most important phase. **Designing** an approach that works best for you requires significant effort up front, but offers a great payout later in the form of **meaningful incident response and overall improved security**.

Source: Crafting the InfoSec Playbook, 4 A Data-Centric Approach to Security Monitoring (bold by me) by Jeff Bollinger, Brandon Enright, and Matthew Valites ISBN: 9781491949405

- I recommend pre-filtering, see what noise your devices *would send*
"Collecting only relevant data can have a direct impact on reducing costs as well."
- Same is recommended in CIP page 50: Just the Facts
- Normalization – "Data normalization for the purposes of log mining is the process by which a portion, or field, of a log event is transformed into its canonical form."

Metadata – enrichment



Metadata	Metacategory: Data
The DNS lookup occurred at a certain time.	Timestamp: 278621182
The internal host sent a DNS PTR request.	Network protocol: DNS PTR
The internal host had a hostname.	Location: Desktop subnet
	Source IP Address: 1.1.1.2
	Hostname: windowspc22.company.com
The internal host resolved an external host.	Location: External
	Destination IP Address: 255.123.215.3
	Hostname: dgf7adfnkjhh.com
The external host was hosted by a dynamic DNS provider.	Network: Shady DDnS Provider Inc.
	ASN: SHADY232
	Reputation: Historically risky network
The remote hostname appeared randomly generated.	Hostname: dgf7adfnkjhh.com
	Category: Unusual, nonlinguistic

Source: picture from *Crafting the InfoSec Playbook*, CIP

Metadata + Context

Summary, what to log



CIP 7 Tools of the Trade, need to know NetFlow, DNS, and HTTP proxy logs in the real-world

- Defense in Depth – we will never catch everything
- Log Management: The Security Event Data Warehouse
- Intrusion Detection Isn't Dead
- DNS, the One True King – Logging and analyzing DNS transactions, Blocking DNS requests or responses
- HTTP Is the Platform: Web Proxies – Web proxies allow you to solve additional security problems
- Rolling Packet Capture – In a perfect world, we would have full packet capture everywhere

Detection Capabilities



Security incidents happen, but what happens. One of the actions to reduce impact of incidents are done in preparing for incidents.

Preparation for an attack, establish procedures and mechanisms for detecting and responding to attacks

Preparation will enable easy **identification** of affected systems, better **containment** which systems are likely to be infected, **eradication** what happened – how to do the **eradication** and **recovery**.

Strategy for implementing identification and detection



We recommend that the following strategy is used for implementing identification and detection.

We have the following recommendations and actions points for logging:

- Enable system logging from servers
- Enable system logging from network devices
- Centralize logging
- Add search facilities and dashboards
- Perform system audits manually or automatically
- Setup notification and notification procedures

Extended Sources



When a basic logging infrastructure is setup, it can be expanded to increase coverage, by adding more sources:

- DNS query logging – will enable multiple cases to be resolved, example malware identification and tracing, when was a malware domain queried, when was the first infection
- Session data from Firewalls, Netflow – traffic patterns can be investigated and both attacks and cases like exfiltration can likely be seen

Hint: Take the sources available first, make a proof-of-concept, expand coverage

Data Analysis Skills



Although we could spend an entire book creating an exhaustive list of skills needed to be a good security data scientist, this chapter covers the following skills/domains that a data scientist will benefit from knowing within information security:

- Domain expertise—Setting and maintaining a purpose to the analysis
- Data management—Being able to prepare, store, and maintain data
- Programming—The glue that connects data to analysis
- Statistics—To learn from the data
- Visualization—Communicating the results effectively

It might be easy to label any one of these skills as the most important, but in reality, the whole is greater than the sum of its parts. Each of these contributes a significant and important piece to the workings of security data science.

Source: *Data-Driven Security: Analysis, Visualization and Dashboards* Jay Jacobs, Bob Rudis

ISBN: 978-1-118-79372-5 February 2014 <https://datadrivensecurity.info/> - short DDS



Baseline Skills

- Threat-Centric Security, NSM, and the NSM Cycle
- TCP/IP Protocols
- Common Application Layer Protocols
- Packet Analysis
- Windows Architecture
- Linux Architecture
- Basic Data Parsing (BASH, Grep, SED, AWK, etc)
- IDS Usage (Snort, Suricata, etc.)
- Indicators of Compromise and IDS Signature Tuning
- Open Source Intelligence Gathering
- Basic Analytic Diagnostic Methods
- Basic Malware Analysis

Source: *Applied Network Security Monitoring Collection, Detection, and Analysis*, Chris Sanders and Jason Smith

Automated packet sniffing tools



The Zeek Network Security Monitor

Zeek – Network Security Monitor <https://zeek.org>

Suricata – open source, mature, fast and robust network threat detection <https://suricata-ids.org/>

ntopng – High-speed web-based traffic analysis <https://www.ntop.org/>

Maltrail – Malicious traffic detection system <https://github.com/stamparm/MalTrail>

Slide included as a reference for Network SECURITY Monitoring

Side note: Zeek Security Monitor handles formats differently



Zeek has files formatted with a header:

```
#fields ts      uid      id.orig_h      id.orig_p      id.resp_h      id.resp_p      proto      trans_id
       rtt      query     qclass    qclass_name    qtype      qtype_name    rcode      rcode_name    AA
       TC       RD       RA        Z           answers    TTLs       rejected
```

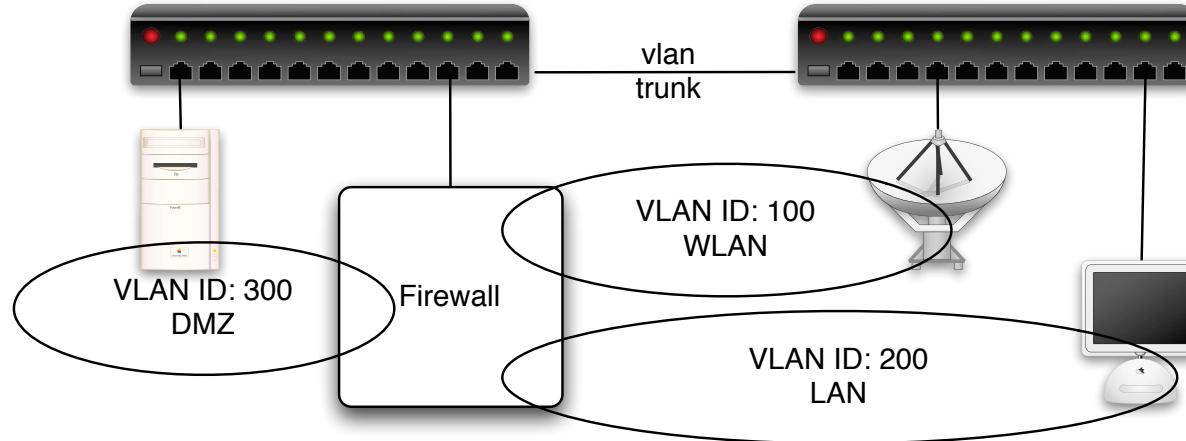
```
1538982372.416180 CD12Dc1SpQm42QW4G3 10.xxx.0.145 57476 10.x.y.141 53 udp 20383
0.045021 www.dr.dk 1 C_INTERNET 1 A 0 NOERROR F F T T 0
    www.dr.dk-v1.edgekey.net,e16198.b.akamaiedge.net,2.17.212.93 60.000000,20409.000000,20.000000 F
```

Note: this show ALL the fields captured and dissected by Zeek, there is a nice utility program bro-cut which can select specific fields:

```
root@NMS-VM:/var/spool/bro/bro# cat dns.log | bro-cut -d ts query answers | grep dr.dk
2018-10-08T09:06:12+0200 www.dr.dk www.dr.dk-v1.edgekey.net,e16198.b.akamaiedge.net,2.17.212.93
```

Can also just use JSON now via Filebeat

IEEE 802.1q



Lets introduce some more advanced things, now that configurations are saved – easy to see changes

Using IEEE 802.1q VLAN tagging on Ethernet frames

Virtual LAN, to pass from one to another, must use a router/firewall

Allows separation/segmentation and protects traffic from many security issues

Port Security – Rogue DHCP servers



Common problem in networks is people connecting devices with DHCPD servers

In general make sure to segment networks

Start to use port security on switches, including DHCP snooping

https://en.wikipedia.org/wiki/DHCP_snooping



Example port security

```
[edit ethernet-switching-options secure-access-port]
set interface ge-0/0/1 mac-limit 4

set interface ge-0/0/2 allowed-mac 00:05:85:3A:82:80
set interface ge-0/0/2 allowed-mac 00:05:85:3A:82:81
set interface ge-0/0/2 allowed-mac 00:05:85:3A:82:83
set interface ge-0/0/2 allowed-mac 00:05:85:3A:82:85
set interface ge-0/0/2 allowed-mac 00:05:85:3A:82:88
set interface ge-0/0/2 mac-limit 4

set interface ge-0/0/1 persistent-learning
set interface ge-0/0/8 dhcp-trusted
set vlan employee-vlan arp-inspection
set vlan employee-vlan examine-dhcp
set vlan employee-vlan mac-move-limit 5
```

Source: Overview of Port Security, Juniper

https://www.juniper.net/documentation/en_US/junos/topics/example/overview-port-security.html

BGP intro



What is BGP Border Gateway Protocol

Dynamic routing protocol, BGPv4 used on whole internet

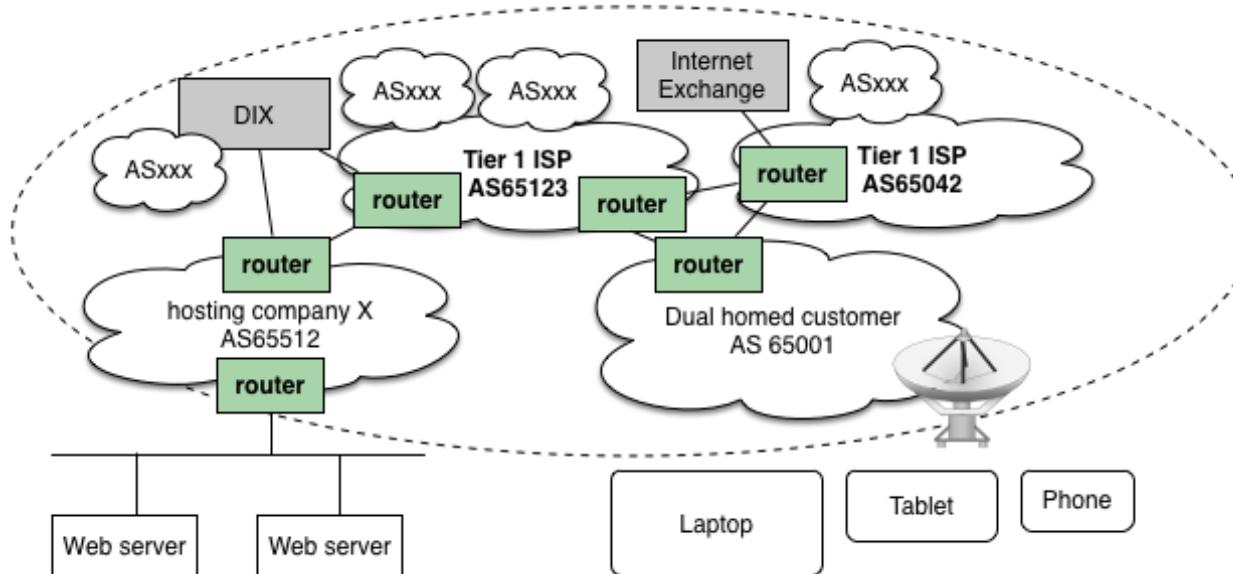
Networks identified using AS numbers ASNs

Autonomous System (AS) can be small or very big, world wide

BGP version 4 RFC-4271 uses TCP connections *peering*

http://en.wikipedia.org/wiki/Border_Gateway_Protocol

Hosting and internet providers



- BGP networks are used for all of the Internet
- New standards like Resource Public Key Infrastructure (RPKI) are underway

Routing and BGP Solutions



- Filtering, ingress / egress:
"reject external packets that claim to be from the local net"
- See also Reverse Path forwarding https://en.wikipedia.org/wiki/Reverse-path_forwarding
- Routers and routing protocols must be more skeptical
Routing filters implemented everywhere, auth on routing protocols OSPF/BGP etc.
- Has been recommended for some years, but not done in all organisations
- BGP routing Resource Public Key Infrastructure RPKI
- BCP38 is RFC2827: *Network Ingress Filtering: Defeating Denial of Service Attacks which employ IP Source Address Spoofing*
<http://www.bcp38.info/>
- *Mutually Agreed Norms for Routing Security*, <https://www.manrs.org/>

Exercise at home – Your lab setup



- Go to GitHub, Find user Kramse, click through security-courses, courses, suricatazeek and download the PDF files for the slides and exercises:

<https://github.com/kramse/security-courses/tree/master/courses/networking/suricatazeek-workshop>

- Get the lab instructions, from

<https://github.com/kramse/kramse-labs/tree/master/suricatazeek>

Questions?



Henrik Kramselund Jereminsen hkj@zecurity.com @kramse  

You are always welcome to send me questions later via email

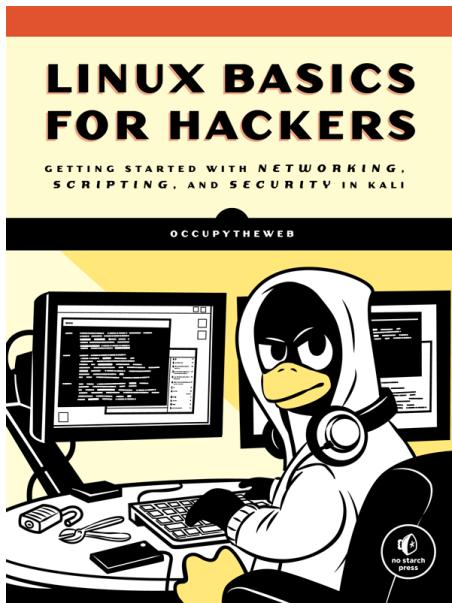
Email: hkj@zecurity.dk Mobile: +45 2026 6000

Recommended further reading



- Campus Network Security: High Level Overview , Network Startup Resource Center https://nsrc.org/workshops/2018/myren-nsrc-cndo/networking/cndo/en/presentations/Campus_Security_Overview.pdf
- Campus Operations Best Current Practice, Network Startup Resource Center https://nsrc.org/workshops/2018/tenet-nsrc-cndo/networking/cndo/en/presentations/Campus_Operations_BCP.pdf
- Mutually Agreed Norms for Routing Security (MANRS) https://www.manrs.org/wp-content/uploads/2018/09/MANRS_PDF_Sep2016.pdf
- RFC2827: Network Ingress Filtering: Defeating Denial of Service Attacks <https://tools.ietf.org/html/rfc2827>

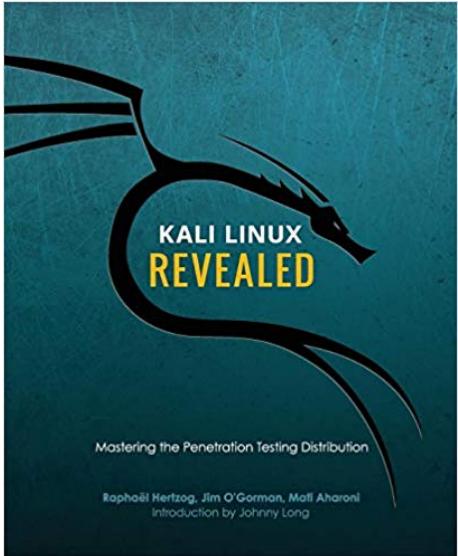
Book: Linux Basics for Hackers (LBhf)



Linux Basics for Hackers Getting Started with Networking, Scripting, and Security in Kali by OccupyTheWeb December 2018, 248 pp. ISBN-13: 9781593278557

<https://nostarch.com/linuxbasicsforhackers> Explains how to use Linux

Book: Kali Linux Revealed (KLR)

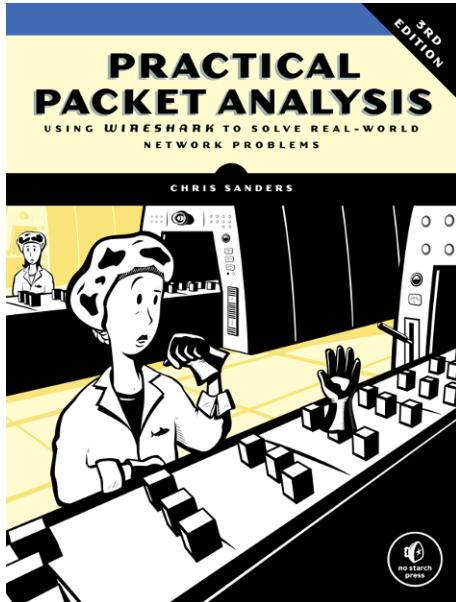


Kali Linux Revealed Mastering the Penetration Testing Distribution

<https://www.kali.org/download-kali-linux-revealed-book/>

Explains how to install Kali Linux

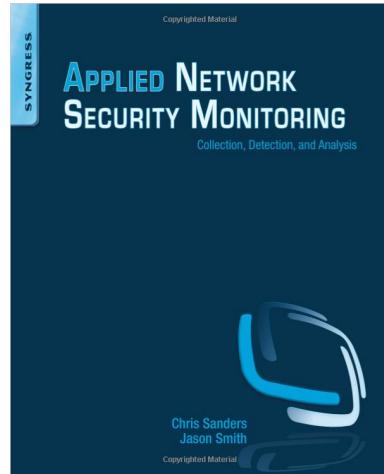
Book: Practical Packet Analysis (PPA)



Practical Packet Analysis, Using Wireshark to Solve Real-World Network Problems by Chris Sanders, 3rd Edition April 2017, 368 pp. ISBN-13: 978-1-59327-802-1

<https://nostarch.com/packetanalysis3>

Book: Applied Network Security Monitoring (ANSM)

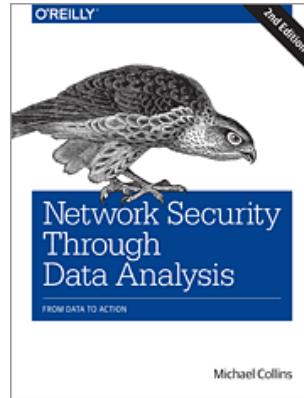


Applied Network Security Monitoring: Collection, Detection, and Analysis 1st Edition

Chris Sanders, Jason Smith eBook ISBN: 9780124172166 Paperback ISBN: 9780124172081 496 pp. Imprint: Syngress, December 2013

<https://www.elsevier.com/books/applied-network-security-monitoring/unknown/978-0-12-417208-1>

Network Security Through Data Analysis



Low page count, but high value! Recommended.

Network Security through Data Analysis, 2nd Edition By Michael S Collins Publisher: O'Reilly Media 2015-05-01:
Second release, 348 Pages

New Release Date: August 2017